

Assignment 5: Image classification

The code in this live script was created based on the starter code. All unattributed text and code should be assumed to be from the starter code. All mistakes and poorly written pieces of code are mine.

All experiments use a subset of Kaggle's "dogs vs. cats" training dataset (<https://www.kaggle.com/c/dogs-vs-cats/data?select=train.zip>).

Experiments are titled as follows:

- Base network name (all networks are customized for transfer learning)
- Test/validation partition
- Optimizer
- Data augmentation (yes/no)

For example, "AlexNet, 70/30, sgdm, no augmentation" is an experiment that uses AlexNet with transfer learning, 70% of the dataset is used for training (and 30% for validation), with the SGDM optimizer, and no data augmentation.

Experiment 1 - AlexNet, 70/30, sgdm, no augmentation

Create the dataset.

```
[trainingSet1, validationSet1] = getDataset(0.7, 'AlexNet');
```


```
Dataset distribution before rebalancing: cat=500, dog=500  
Dataset distribution after rebalancing: cat=500, dog=500  
Training set distribution: cat=350, dog=350  
Validation set distribution: cat=150, dog=150
```

Create the model.

```
model1 = alexNetTransferLearning();
```

Train the model.

```
options1 = trainingOptions('sgdm', ...  
    'MiniBatchSize',10, ...  
    'MaxEpochs',6, ...  
    'InitialLearnRate',1e-4, ...  
    'Shuffle','every-epoch', ...  
    'ValidationData',validationSet1, ...  
    'ValidationFrequency',3, ...  
    'Verbose',false, ...  
    'Plots','training-progress');  
  
[YPred1,scores1] = trainEvaluate(model1,options1,trainingSet1,validationSet1);
```



Accuracy and confusion matrix.

```
showPredictionResults(validationSet1, YPred1)
```

282 of 300 images were predicted correctly

The validation accuracy is: 94.00 %

Confusion matrix

True Class	cat	142	8
	dog	10	140

		93.4%	94.6%
		6.6%	5.4%
		cat	dog
		Predicted Class	

Some correctly predicted images.


```
showImages(validationSet1, YPred1, scores1, true)
```

Correct predictions

Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000




Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=0.802494, dog=0.197506



Some incorrectly predicted images.

```
showImages(validationSet1, YPred1, scores1, false)
```

Incorrect predictions

Prediction: dog, scores: cat=0.077625, dog=0.922375



Prediction: dog, scores: cat=0.000325, dog=0.999675



Prediction: dog, scores: cat=0.456494, dog=0.543506



Prediction: dog, scores: cat=0.241180, dog=0.758820



Prediction: dog, scores: cat=0.000077, dog=0.999923



Experiment 2 - AlexNet, 80/20, sgdm, no agumentation

Create the dataset.

```
[trainingSet2, validationSet2] = getDataset(0.8, 'AlexNet');
```

Dataset distribution before rebalancing: cat=500, dog=500

Dataset distribution after rebalancing: cat=500, dog=500

Training set distribution: cat=400, dog=400

Validation set distribution: cat=100, dog=100

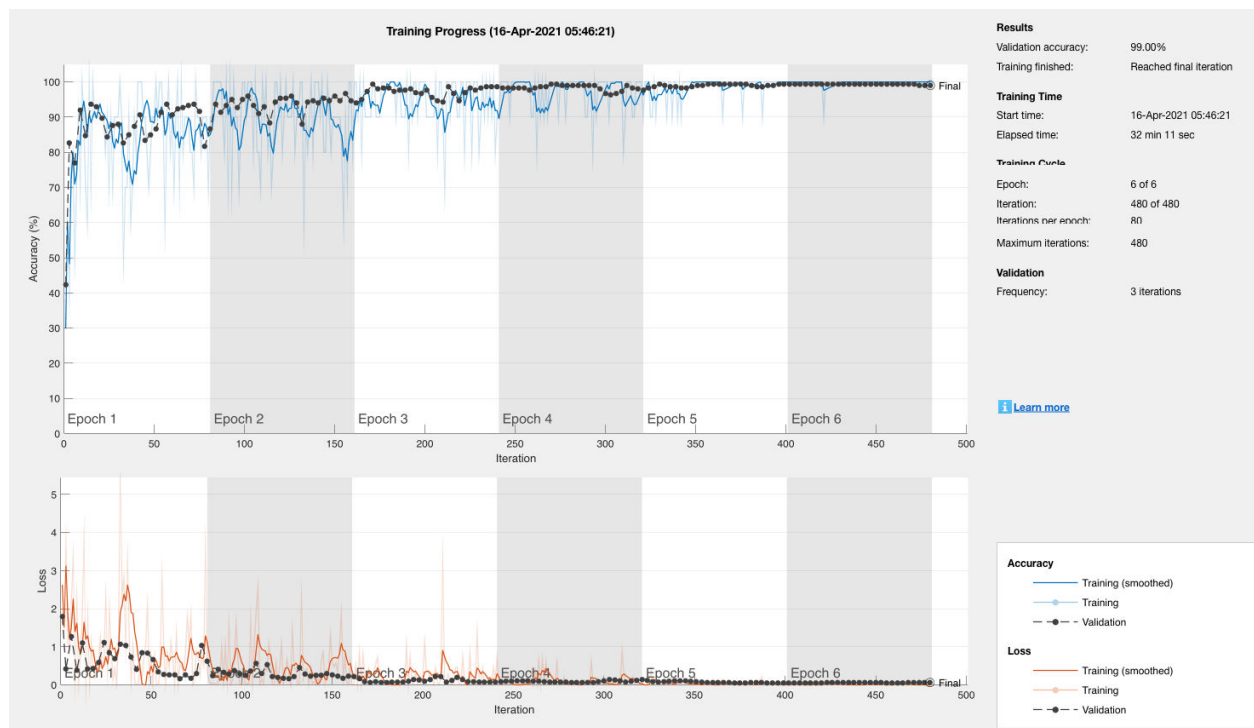
Create the model.

```
model2 = alexNextTransferLearning();
```

Train the model.

NOTE that we are using the same training options from the first experiment.

```
[YPred2,scores2] = trainEvaluate(model2,options1,trainingSet2,validationSet2);
```



Accuracy and confusion matrix.

```
showPredictionResults(validationSet2, YPred2)
```

185 of 200 images were predicted correctly
The validation accuracy is: 92.50 %

Confusion matrix


True Class	cat	96	4
	dog	11	89
		89.7%	95.7%
		10.3%	4.3%
		cat	dog
		Predicted Class	

Some correctly predicted images.


```
showImages(validationSet2, YPred2, scores2, true)
```

Correct predictions

Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=0.999995, dog=0.000005



Some incorrectly predicted images.

```
showImages(validationSet2, YPred2, scores2, false)
```

Incorrect predictions


Prediction: dog, scores: cat=0.006176, dog=0.993824



Prediction: dog, scores: cat=0.000033, dog=0.999967



Prediction: dog, scores: cat=0.003065, dog=0.996935



Prediction: dog, scores: cat=0.000535, dog=0.999465



Prediction: cat, scores: cat=0.734694, dog=0.265306



Experiment 3 - AlexNet, 80/20, rmsprop, no agumentation

Create the dataset.

```
[trainingSet3, validationSet3] = getDataset(0.8, 'AlexNet');
```


Dataset distribution before rebalancing: cat=500, dog=500
Dataset distribution after rebalancing: cat=500, dog=500
Training set distribution: cat=400, dog=400
Validation set distribution: cat=100, dog=100

Create the model.

```
model3 = alexNetTransferLearning();
```

Train the model.

```
options3 = trainingOptions('rmsprop', ...  
    'MiniBatchSize',10, ...  
    'MaxEpochs',6, ...  
    'InitialLearnRate',1e-4, ...  
    'Shuffle','every-epoch', ...  
    'ValidationData',validationSet3, ...  
    'ValidationFrequency',3, ...  
    'Verbose',false, ...  
    'Plots','training-progress');  
  
[YPred3,scores3] = trainEvaluate(model3,options3,trainingSet3,validationSet3);
```



Accuracy and confusion matrix.

```
showPredictionResults(validationSet3, YPred3)
```

178 of 200 images were predicted correctly

The validation accuracy is: 89.00 %

Confusion matrix

True Class	cat	87	13
	dog	9	91
		90.6%	87.5%
		9.4%	12.5%
		cat	dog
		Predicted Class	

Some correctly predicted images.

```
showImages(validationSet3, YPred3, scores3, true)
```

Correct predictions

Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=0.992314, dog=0.007686



Prediction: cat, scores: cat=0.723763, dog=0.276237



Some incorrectly predicted images.


```
showImages(validationSet3, YPred3, scores3, false)
```

Incorrect predictions

Prediction: dog, scores: cat=0.000000, dog=1.000000




Prediction: dog, scores: cat=0.006982, dog=0.993018



Prediction: dog, scores: cat=0.000000, dog=1.000000



Prediction: dog, scores: cat=0.000000, dog=1.000000



Prediction: dog, scores: cat=0.113243, dog=0.886757



Experiment 4 - GoogLeNet, 80/20, rmsprop, no agumentation

Create the dataset.

```
[trainingSet4, validationSet4] = getDataset(0.8, 'GoogLeNet');
```

Dataset distribution before rebalancing: cat=500, dog=500

Dataset distribution after rebalancing: cat=500, dog=500

Training set distribution: cat=400, dog=400

Validation set distribution: cat=100, dog=100

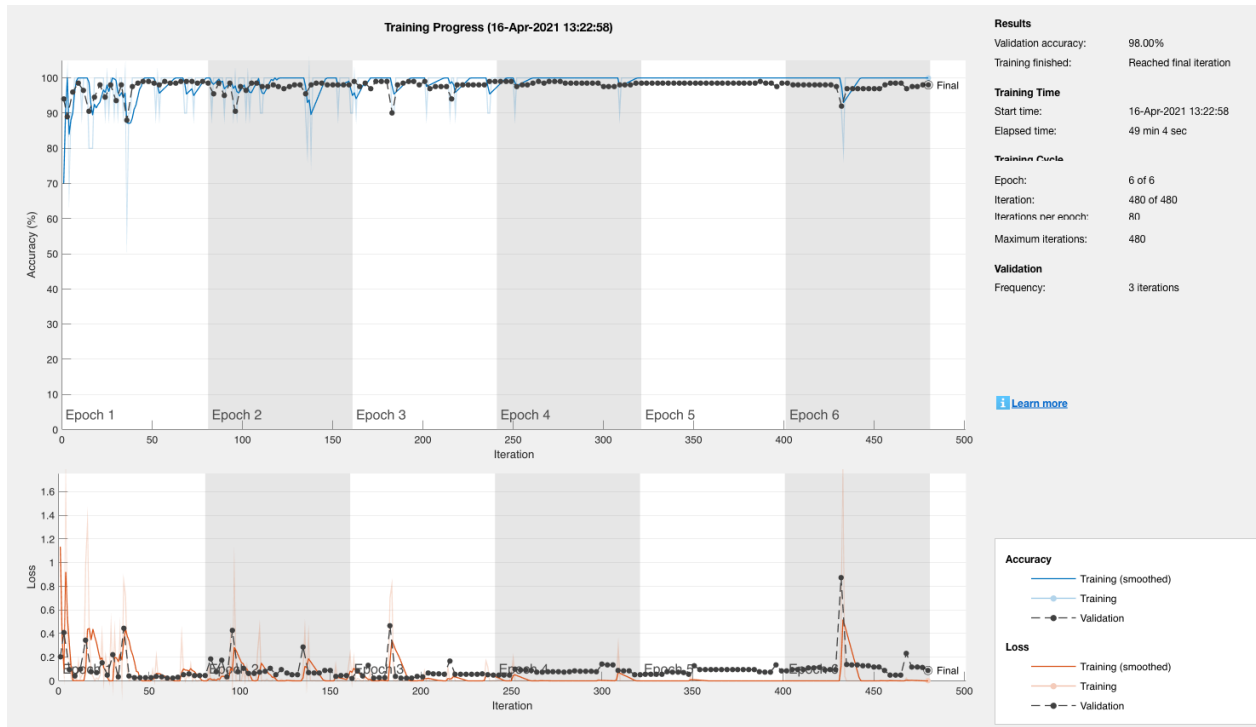
Create the model.

```
model4 = GoogLeNetTansferLearning();
```

Train the model.

```
options4 = trainingOptions('rmsprop', ...  
    'MiniBatchSize',10, ...  
    'MaxEpochs',6, ...  
    'InitialLearnRate',1e-4, ...  
    'ValidationData',validationSet4, ...  
    'ValidationFrequency',3, ...  
    'ValidationPatience',Inf, ...  
    'Verbose',false, ...  
    'Plots','training-progress');
```

```
[YPred4,scores4] = trainEvaluate(model4,options4,trainingSet4,validationSet4);
```



Accuracy and confusion matrix.

```
showPredictionResults(validationSet4, YPred4)
```

196 of 200 images were predicted correctly
The validation accuracy is: 98.00 %

Confusion matrix

True Class	cat	99	1
	dog	3	97
		97.1%	99.0%
		2.9%	1.0%
		cat	dog
		Predicted Class	

Some correctly predicted images.

```
showImages(validationSet4, YPred4, scores4, true)
```

Correct predictions
Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Some incorrectly predicted images.


```
showImages(validationSet4, YPred4, scores4, false)
```

Incorrect predictions

Prediction: dog, scores: cat=0.339965, dog=0.660035




Prediction: cat, scores: cat=0.999991, dog=0.000009



Prediction: cat, scores: cat=0.933207, dog=0.066793



Prediction: cat, scores: cat=0.640845, dog=0.359155



Experiment 5 - GoogLeNet, 80/20, rmsprop, with augmentation

Create the model.

```
model5 = GoogLeNetTransferLearning();
```


Create an augmented dataset.

```
[trainingSet5, validationSet5] = getDataset(0.8, 'GoogLeNet');
```

Dataset distribution before rebalancing: cat=500, dog=500

Dataset distribution after rebalancing: cat=500, dog=500

Training set distribution: cat=400, dog=400

Validation set distribution: cat=100, dog=100

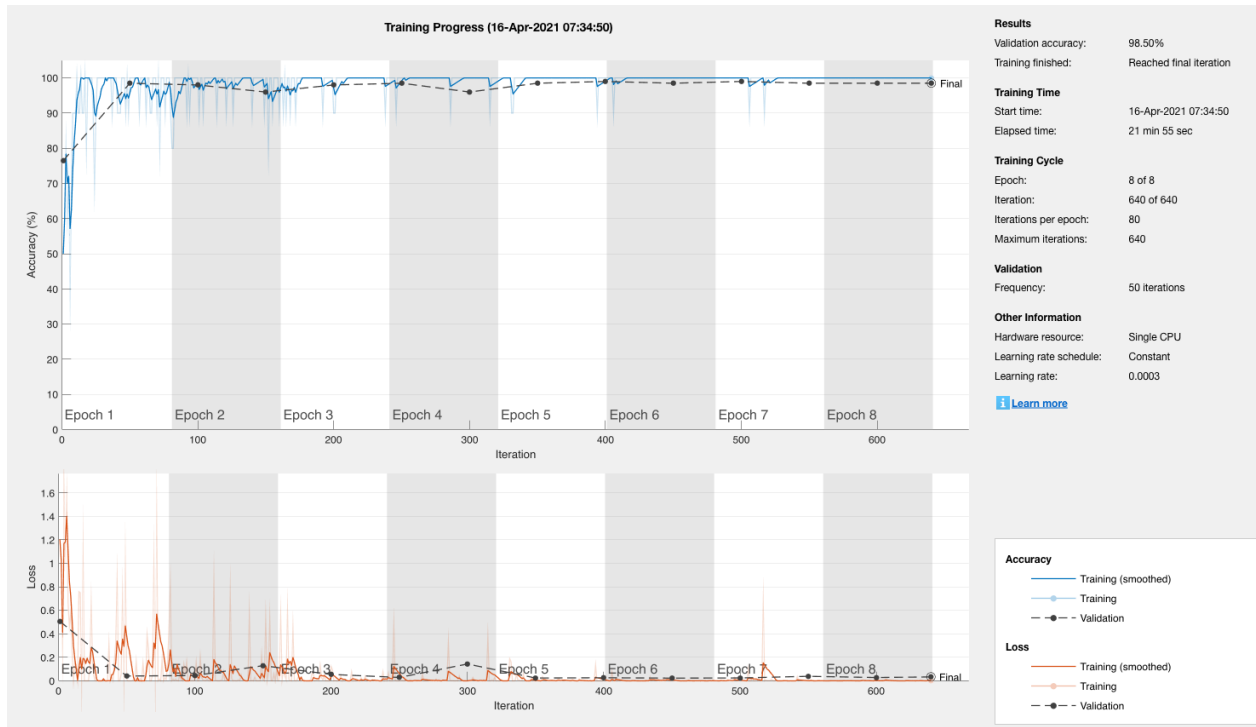
```
% In our case, we shall use an augmented image datastore to randomly flip
% the training images along the vertical axis and randomly translate them
% up to 30 pixels and scale them up to 10% horizontally and vertically.
pixelRange = [-30 30];
scaleRange = [0.9 1.1];
imageAugmenter = imageDataAugmenter( ...
    'RandXReflection',true, ...
    'RandXTranslation',pixelRange, ...
    'RandYTranslation',pixelRange, ...
    'RandXScale',scaleRange, ...
    'RandYScale',scaleRange);

inputSize = model5.Layers(1).InputSize;
augimdsTrain = augmentedImageDatastore(inputSize(1:2),trainingSet5, ...
    'DataAugmentation',imageAugmenter);
augimdsValidation = augmentedImageDatastore(inputSize(1:2),validationSet5);
```

Train the model.

```
miniBatchSize = 10;
options5 = trainingOptions('sgdm', ...
    'MiniBatchSize',miniBatchSize, ...
    'MaxEpochs',8, ...
    'InitialLearnRate',3e-4, ...
    'ValidationData',augimdsValidation, ...
    'Verbose',false, ...
    'Plots','training-progress');

trainedModel5 = trainNetwork(augimdsTrain,model5,options5);
```



Predict scores on the validation set and calculate accuracy.

```
[YPred5,scores5] = classify(trainedModel5,augimdsValidation);
```

Accuracy and confusion matrix.

```
showPredictionResults(validationSet5, YPred5)
```

197 of 200 images were predicted correctly
The validation accuracy is: 98.50 %

Confusion matrix

True Class	cat	97	3
	dog		100
		100.0%	97.1%
			2.9%
		cat	dog
		Predicted Class	

Some correctly predicted images.

```
showImages(validationSet5, YPred5, scores5, true)
```

Correct predictions
Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Prediction: cat, scores: cat=1.000000, dog=0.000000



Some incorrectly predicted images.

```
showImages(validationSet5, YPred5, scores5, false)
```

Incorrect predictions

Prediction: dog, scores: cat=0.206454, dog=0.793546



Prediction: dog, scores: cat=0.239825, dog=0.760175



Prediction: dog, scores: cat=0.057522, dog=0.942478



Auxiliary functions

Dataset

Create a dataset backed by an image store. The function assumes there are only two categories and that the directory structure is correctly split into these categories (the name of the directories must match the name of the categories for the `imageDatastore` to work correctly).

```
function [trainingSet, validationSet] = getDataset(trainPercentage, network)

    CATEGORIES = {'cat', 'dog'};
    PATH = './data/kaggle';

    function showStats(text, imds)
        tempTbl = countEachLabel(imds);
```

```

        fprintf('%s: %s=%d, %s=%d\n', text, ...
            tempTbl(1,1).Label, tempTbl(1,2).Count, ...
            tempTbl(2,1).Label, tempTbl(2,2).Count)
    end

    imds = imageDatastore(fullfile(PATH, CATEGORIES), 'LabelSource', 'foldernames');
    showStats('Dataset distribution before rebalancing', imds)

    % Ensure the categories are balanced by using the smallest overlap set (may not be
    % needed for some datasets, but protects in case we use an unbalanced dataset
    % without realizing it).
    tbl = countEachLabel(imds);
    minSetCount = min(tbl{:,2});
    % Trim the set.
    imds = splitEachLabel(imds, minSetCount, 'randomize');
    % Each set must now have exactly the same number of images.
    showStats('Dataset distribution after rebalancing', imds)

    % AlexNet can only process RGB images that are 227-by-227
    % To avoid re-saving all the images to this format, setup the |imds| read function
    % |imds.ReadFcn|, to pre-process images on-the-fly. The |imds.ReadFcn| is called
    % every time an image is read from the |ImageDatastore|
    if isequal(network, 'AlexNet')
        imds.ReadFcn = @(filename) readAndPreprocessImageAlexNet(filename);
    else
        imds.ReadFcn = @(filename) readAndPreprocessImageGoogLeNet(filename);
    end

    [trainingSet, validationSet] = splitEachLabel(imds, trainPercentage, 'randomized')
    showStats('Training set distribution', trainingSet)
    showStats('Validation set distribution', validationSet)
end

```

Model

Create an AlexNet model ready for transfer learning.

```

function model = alexNextTransferLearning()
    originalModel = alexnet;

    layersTransfer = originalModel.Layers(1:end-3);

    NUM_CLASSES = 2;

    model = [
        layersTransfer
        fullyConnectedLayer(NUM_CLASSES, 'WeightLearnRateFactor', 20, 'BiasLearnRateFactor', 20)
        softmaxLayer
        classificationLayer];
end

```

Create a GoogLeNet model ready for transfer learning. Based on <https://www.mathworks.com/help/releases/R2018a/nnet/examples/transfer-learning-using-googlenet.html>.

```
function model = GoogLeNetTransferLearning()
    originalModel = googlenet;

    lgraph = layerGraph(originalModel);
    lgraph = removeLayers(lgraph, {'loss3-classifier','prob','output'});

    NUM_CLASSES = 2;

    newLayers = [
        fullyConnectedLayer(NUM_CLASSES,'Name','fc','WeightLearnRateFactor',10,'BiasLearnRateFactor',10)
        softmaxLayer('Name','softmax')
        classificationLayer('Name','classoutput')];
    lgraph = addLayers(lgraph,newLayers);

    model = connectLayers(lgraph,'pool5-drop_7x7_s1','fc');
end
```

Train and evaluate a model.

```
function [YPred,scores] = trainEvaluate(model,options,trainingSet,validationSet)
    trainedModel = trainNetwork(trainingSet,model,options);
    [YPred,scores] = classify(trainedModel,validationSet);
end
```

Show the prediction results.

```
function showPredictionResults(validationSet, YPredicted)
    YValidation = validationSet.Labels;

    correct = YPredicted == YValidation;
    fprintf("\n\n%d of %d images were predicted correctly\n", ...
        sum(correct), numel(YValidation));
    fprintf("The validation accuracy is: %.2f %%\n",mean(correct) * 100);

    fprintf("\nConfusion matrix\n")
    confusionchart(YValidation,YPredicted, ...
        "ColumnSummary","column-normalized")
end
```

Show correctly/incorrectly images.

```
function showImages(validationSet, YPredicted, scores, showCorrect)
    if showCorrect
        fprintf("Correct predictions\n")
    else
        fprintf("Incorrect predictions\n")
    end
end
```

```

        fprintf("Incorrect predictions\n")
    end

    NUM_IMAGES = 5;
    YValidation = validationSet.Labels;

    shownImages = 0;
    for i = 1:numel(YValidation)
        if showCorrect
            show = (YValidation(i) == YPredicted(i));
        else
            show = (YValidation(i) ~= YPredicted(i));
        end

        if show
            shownImages = shownImages + 1;
            if shownImages > NUM_IMAGES
                break
            end

            fprintf("Prediction: %s, scores: cat=%f, dog=%f\n", ...
                YPredicted(i), scores(i,1), scores(i,2))
            figure
            img = readimage(validationSet,i);
            imshow(img)
            drawnow
        end
    end
end
end

```